

6.034 Quiz 4

1 December 2010

Name	Uldis Sturms
email	uldis.sturms@gmail.com

Circle your TA and recitation time, so that we can more easily enter your score in our records and return your quiz to you promptly.

TAs
Martin Couturier
Kenny Donahue
Bobby Keys
Gleb Kuznetsov
Kendra Pugh
Mark Seifter
Yuan Shen

Thu	
Time	Instructor
1-2	Bob Berwick
2-3	Bob Berwick
3-4	Bob Berwick

Fri	
Time	Instructor
1-2	Randall Davis
2-3	Randall Davis
3-4	Randall Davis

Problem number	Maximum	Score	Grader
1	50		
2	50		
Total	100		


There are 9 pages in this quiz, including this one, but not including blank pages and tear-off sheets. Tear-off sheets are provided at the end with duplicate drawings and data. **As always, open book, open notes, open just about everything, including a calculator, but no computers.**

Problem 1: SVMs (50 points)

After reading too much Lord of the Rings, you wake up to find yourself in Middle Earth. You decide the most relevant thing to do is to classify the different races around you.

Part A: Distinguishing Dwarves from Humans (38 points)

You meet 2 dwarves (+) and 3 humans (-) and realize that they have distinguishing features: beard length (x) and height (y). You plot these data points on a grid. Being an expert in SVMs, you decide to start off on an epic journey of classification.



A1 (7 points)

Draw the **decision boundary** on the graph above and **clearly label positive and negative gutters**, and **circle all support vectors**.

What is the width of the road/margin?

$y = x - 2$ (decision boundary)
 $y = -(w_1/w_2)x - b/w_2$, $w_1/w_2 = -1$, $b/w_2 = 2$
 $1/|w| = \sqrt{2}$
 $|w| = (\sqrt{2})/2$

A2 (12 points)

Compute \vec{w} and b in the decision boundary $h(\vec{u}) = \vec{w} \cdot \vec{u} + b \geq 0$ for the SVM solution to part A1.

Show your work here.

$$\sqrt{(w_1^2 + w_2^2)} = (\sqrt{2})/2, w_1\sqrt{2} = -(\sqrt{2})/2, w_1 = -0.5, w_2 = 0.5, b = 1$$

$$\vec{w} = 0.5$$

$$\mathbf{b} = 1$$

A2 (10 points)

Calculate the weights (alphas) of each data point.

Show your work here.

$$-0.5 = 3\alpha - 5\alpha$$

$$0.5 = 3\alpha - 1\alpha$$

$$\alpha = 0.25$$

$$\alpha_1 = 0.25$$

$$\alpha_2 = 0.25$$

$$\alpha_3 = 0.25$$

$$\alpha_4 = 0.25$$

$$\alpha_5 = 0.25$$

A3 (9 points)

What will be the alpha of a new negative point 6 placed at (0, 6)?

0.25

What will be the alpha of a new negative point 6 placed at (0,0)?

0.25

Supposed we moved point 3 to (4, 2), how will the **magnitude of the alpha 3** change?

Circle one:

Larger

Smaller

Same


Part B: Distinguishing Kernels (12 points)

Back in his lab, Gandolf has been hacking on some kernels in preparation for greater classification adventures. For each of the following, indicate **YES** or **NO** whether the kernel can be used to *perfectly* classify the test points, and if **YES** *sketch* the decision boundaries and gutters (the street) such a classifier might produce and *circle* which data points are support vectors. Note that because of symmetry, more than one answer may be possible for one or more cases.

$$K(\vec{u}, \vec{v}) = \vec{u} \cdot \vec{v}$$

YES


NO



$$K(\vec{u}, \vec{v}) = (\vec{u} \cdot \vec{v} + 1)^2$$

YES


NO



$$K(\vec{u}, \vec{v}) = e^{-(|\vec{u} - \vec{v}|)^2 / 2}$$

YES

NO



Problem 2: Boosting (50 points)

After wearing Sauron's ring for several months, Frodo is rapidly losing his sanity. He fears that the ring will interfere with his better judgement and betray him to an enemy. To ensure that he doesn't put his trust into enemy hands, he flees Middle Earth in search of a way to classify his enemies from his friends. In his travels he had heard rumors of the magic of Artificial Intelligence and has decided to hire you to build him a classifier, which will correctly differentiate between his friends and his enemies. Below is all of the information Frodo remembers about the people back in Middle Earth.

ID	Name	Friend	Species	Has Magic	Part of the Fellowship	Has/Had a ring of power	Length of hair (feet)
1	Gandalf	Yes	Wizard	Yes	Yes	No	2
2	Sarumon	No	Wizard	Yes	No	No	2.5
3	Sauron	No	Wizard	Yes	No	Yes	0
4	Legolas	Yes	Elf	Yes	Yes	No	2
5	Tree-Beard	Yes	Ent	No	No	No	0
6	Sam	Yes	Hobbit	No	Yes	No	0.25
7	Elrond	Yes	Elf	Yes	No	Yes	2
8	Gollum	No	Hobbit	No	No	Yes	1
9	Aragorn	Yes	Man	No	Yes	No	0.75
10	Witch-king of Angmar	No	Man	Yes	No	Yes	2.5

Part A: Picking Classifiers (10 points)

A1 (6 points)

The data has a high dimensionality and so rather than trying to learn an SVM in a high dimension space you think it would be a smart approach to come up with a series of 1 dimensional stubs that can be used to construct a boosting classifier. Fill in the classifier table below. Each of the different classifiers are given a unique ID and a test returns +1 (friend) if true and -1 (enemy) if false.

Classifier	Test	Misclassified
A	Species is a Wizard	2, 3, 4, 5, 6, 7, 9
B	Species is an Elf	1, 5, 6, 9
C	Species is not a Man	2, 3, 8, 9
D	Does not have magic	1, 4, 7, 8
E	Is not part of the Fellowship	1, 2, 3, 4, 6, 8, 9, 10
F	Has never owned a ring of power	2, 7
G	Hair \leq 1ft	1, 2, 3, 4, 7, 8, 10
H	Hair \leq 2 ft	3, 8
I	Friend	2, 3, 8, 10
J	Enemy	1, 4, 5, 6, 7, 9

A2 (4 points)

Looking at the results of your current classifiers, you quickly see two more good weak classifiers (make fewer than 4 errors). What are they?

Classifier	Test	Misclassified
K	Is part of the Fellowship	5, 7
L	Species is not a Wizard	1, 8, 10

Part B: Build a Strong Classifier (30 points)

B1 (25 points)

You realize that many of your tests are redundant and decide to move forward using only these four classifiers: $\{\mathbf{B}, \mathbf{D}, \mathbf{F}, \mathbf{I}\}$. Run the Boosting algorithm on the dataset with these four classifiers. Fill in the weights, classifiers, errors and alphas for three rounds of boosting. In case of ties, favor classifiers that come first alphabetically.

	Round 1		Round 2		Round 3	
w1	0.1	$h_1 = \mathbf{F}$	1/16	$h_2 = \mathbf{B}$	3/24	$h_3 = \mathbf{I}$
w2	0.1	Err = 0.2	4/16	Err = 0.25	4/24	Err = 7/24 ≈ 0.3
w3	0.1	$\alpha = 0.5 \cdot \ln(4)$ $\cong 0.7$	1/16	$\alpha = 0.5 \cdot \ln(3)$ $\cong 0.55$	1/24	$\alpha = 0.5 \cdot \ln(2.33)$ $\cong 0.425$
w4	0.1		1/16		1/24	
w5	0.1		1/16		3/24	
w6	0.1		1/16		3/24	
w7	0.1		4/16		4/24	
w8	0.1		1/16		1/24	
w9	0.1		1/16		3/24	
w10	0.1		1/16		1/24	
Err(B)	0.4		1/4		1/2	
Err(D)	0.4		7/16		3/8	
Err(F)	0.2		1/2		1/3	
Err(I)	0.4		7/16		7/24	

B2 (5 points)

What is the resulting classifier that you obtain after three rounds of Boosting?

$$H(x) = \text{Sign}(0.7 \cdot F(x) + 0.55 \cdot B(x) + 0.425 \cdot I(x))$$

Part C: Boost by Inspection (10 points)


As you become frustrated that you must have picked the wrong subset of classifiers to work with, one of the 6.034 TA's, Martin, happens to walk by and sees your answer to part A1. He reminds you why the boosting algorithm works and then tells you that there is no reason to actually run boosting on this dataset. A boosted classifier of the form:

$$H(x) = \text{Sign}[h_1(x) + h_2(x) + h_3(x)]$$

can be found which solves the problem. What three classifiers $\{h_1, h_2, h_3\}$ is Martin referring to, and why is the resulting $H(x)$ guaranteed to classify all of the points correctly?


$h_1 = B$
 $h_2 = F$
 $h_3 = H$

Tear off sheets. You need not hand these in.




$$K(\vec{u}, \vec{v}) = \vec{u} \cdot \vec{v}$$

YES NO




$$K(\vec{u}, \vec{v}) = (\vec{u} \cdot \vec{v} + 1)^2$$

YES NO



$$K(\vec{u}, \vec{v}) = e^{-(|\vec{u} - \vec{v}|)^2/2}$$

YES NO



ID	Name	Friend	Species	Has Magic	Part of the Fellowship	Has/Had a ring of power	Length of hair (feet)
1	Gandalf	Yes	Wizard	Yes	Yes	No	2
2	Sarumon	No	Wizard	Yes	No	No	2.5
3	Sauron	No	Wizard	Yes	No	Yes	0
4	Legolas	Yes	Elf	Yes	Yes	No	2
5	Tree-Beard	Yes	Ent	No	No	No	0
6	Sam	Yes	Hobbit	No	Yes	No	0.25
7	Elrond	Yes	Elf	Yes	No	Yes	2
8	Gollum	No	Hobbit	No	No	Yes	1
9	Aragorn	Yes	Man	No	Yes	No	0.75
10	Witch-king of Angmar	No	Man	Yes	No	Yes	2.5

MIT OpenCourseWare
<http://ocw.mit.edu>

6.034 Artificial Intelligence
Fall 2010

For information about citing these materials or our Terms of Use, visit: <http://ocw.mit.edu/terms>.